

# Continuous-Time Radar-Inertial Odometry with B-Spline Sliding-Window Optimization

Timo Weiss and Lukas Pries

**Abstract**—We present a radar-inertial odometry (RIO) system for 6-DoF state estimation on a quadrotor using a single-chip 60 GHz mmWave radar (10–60 points per frame) and a 1 kHz IMU, evaluated through aggressive racing and attitude-tumbling flight up to  $10 \text{ rad s}^{-1}$ : a regime existing RIO evaluations do not cover, where the radar model itself degrades. The trajectory is a continuous-time quintic position B-spline and a cumulative SO(3) orientation B-spline, optimized by a fixed-lag smoother with Schur complement marginalization at a 0.3 s stride over a 3 s window; a sensor-only pipeline initializes it without external pose reference. Doppler measures velocity directly, so we evaluate RIO as a *velocity and orientation* source, reporting position as drift per distance. A Markov-blanket factor-selection rule makes the marginalization a true marginal, and a single dynamics-adaptive measurement-weighting law covers hover to  $10 \text{ rad s}^{-1}$ , validated on held-out racing flights. At the live leading edge (the newest window estimate, before later windows refine it), velocity and orientation RMSE stay below  $0.5 \text{ m s}^{-1}$  and  $3^\circ$  across racing (position drift under 1 %) and degrade gracefully to  $6.4^\circ$  orientation (2.9 % drift) under tumbling; the output orientation covariance is approximately calibrated (NEES) on the racing flights, and orientation holds within  $0.3^\circ$  under a realistic magnetometer rather than idealized yaw. We also document negative results: a diagnostic for marginalization inconsistency, and that the tumbling elevation bias is a flight-regime proxy, not an antenna constant.

**Index Terms**—radar-inertial odometry, continuous-time estimation, B-spline, sliding window, marginalization, mmWave radar, Doppler velocity

## I. INTRODUCTION

Visual odometry and light detection and ranging (LiDAR)-based simultaneous localization and mapping (SLAM) are well-established, but both degrade in dusty, foggy, featureless, or textureless environments. Millimeter-wave (mmWave) frequency-modulated continuous-wave (FMCW) radar is inherently robust to these conditions: it penetrates smoke and dust, operates in the dark, and measures Doppler velocity directly without feature matching. These properties make it attractive for aggressive flight in unstructured environments.

Combining a mmWave radar with an inertial measurement unit (IMU) (radar-inertial odometry (RIO)) for ego-motion estimation has been explored with discrete-time extended Kalman filter (EKF) and graph-optimization formulations [1], [2]. However, the asynchronous nature of radar-IMU data (radar at  $\sim 11 \text{ Hz}$ , IMU at  $\sim 1 \text{ kHz}$ ) motivates a continuous-time representation where the trajectory is a smooth function

evaluated at any sensor timestamp without interpolation approximation.

**Research question.** This work asks, end to end: *can a single-chip mmWave radar, fused with an IMU, provide usable 6-DoF state estimation, and if so, what does it take?* The starting point is what the sensor can and cannot do: per-point Doppler yields direct ego-velocity, but at 10–60 points per frame there are no bearing-stable features for position fixing, and yaw is a gauge freedom of Doppler+IMU. The system is therefore a *velocity and orientation* source whose position estimate drifts with distance, by construction rather than by defect. What it takes to turn that capability into accurate estimates, and where it ultimately fails, is the subject of the rest of this paper.

We contribute:

- *Single-chip radar under aggressive and tumbling flight*, characterizing how the radar measurement model breaks in this regime: rate-dependent intra-frame smearing, and the finding that the apparent elevation bias under tumbling is a flight-regime error proxy, not a physical antenna constant. Backflips at  $10 \text{ rad s}^{-1}$  on a single-chip sensor are, to our knowledge, unexplored.
- A *consistency-analyzed* fixed-lag smoother: a factor-set selection rule exploiting B-spline local support to make the Schur marginal exactly closed, distinct from and complementary to first-estimates Jacobian (FEJ), eliminating prior-weight tuning, with a real-time operating point per regime at minimal accuracy cost (0.3–0.5 s stride on a laptop).
- One *weighting law* from hover to  $10 \text{ rad s}^{-1}$  (with per-regime spline grids and two tumbling rescue terms), the velocity–orientation trade a measured Pareto curve, validated on *held-out racing flights* (tumbling awaits the repeated flips of Section VII-A).
- Honest negative results: a diagnostic symptom pattern for marginalization inconsistency; a representation-level study that *disproves* the adaptive-knot-density hypothesis for this regime; ground-truth limits during aggressive maneuvers; and a normalized estimation error squared (NEES) check confirming the reported orientation covariance is consistent on racing flights.

**Data and code.** The complete system is released for independent reproduction. The datasets and the C++ Ceres solver with Python bindings are available at <https://github.com/Dr1zz3l/spline-rio>.

## II. RELATED WORK

A recent survey [3] catalogues the rapidly growing use of mmWave radar in robotics across motion estimation, localiza-

Manuscript received Month XX, 2026; revised Month XX, 2026.

The authors are with the Department of Informatics, Technical University of Munich (TUM), 80333 Munich, Germany (e-mail: timo.weiss@tum.de; lukas.pries@tum.de).

Code and datasets are available at <https://github.com/Dr1zz3l/spline-rio>.

tion, and mapping; we focus on the radar-inertial odometry thread most relevant to single-chip aerial fusion.

**Radar ego-velocity estimation.** Doppler-based ego-velocity estimation from mmWave radar was demonstrated by Kellner et al. [4] and extended to 3D with weighted least squares by Kramer et al. [2], whose RVE system fuses radar with IMU in a discrete-time EKF. Doer and Trommer [1] present an EKF-based RIO with online extrinsic calibration. Our work uses the same per-frame weighted least squares (WLS) ego-velocity estimator but embeds it in a continuous-time batch/sliding-window optimizer.

**Continuous-time trajectory estimation.** Furgale et al. [5] introduced uniform B-splines for continuous-time visual-inertial calibration. The basalt visual-inertial odometry (VIO) [6] uses cumulative B-splines on Lie groups for stereo-inertial odometry; we adapt this basalt-style cumulative SO(3) B-spline to radar-inertial fusion.

**Marginalization in sliding-window estimators.** OKVIS [7] and VINS-Mono [8] pioneered Schur complement marginalization for visual-inertial systems, propagating information from dropped keyframes as a dense Gaussian prior. We apply the same principle to continuous-time B-spline control points and document the practical challenges of prior scaling when the marginal's orientation and position information differ by ten orders of magnitude.

**Closest related work.** Continuous-time radar-inertial odometry was introduced by Ng et al. [9], who fuse Doppler measurements from *multiple automotive* radars with an IMU as a least-squares problem over a B-spline trajectory, evaluated on passenger vehicles. Burnett et al. [10] use a Gaussian-process motion prior for spinning-radar- and lidar-inertial odometry, and Štironja et al. [11] model the IMU trajectory continuously inside an RIO estimator with online spatio-temporal calibration. Single-chip radar-inertial estimation has itself received recent attention: Huang et al. [12] exploit radar cross-section and Doppler physics to make a sparse single-chip cloud usable in a discrete-time sliding-window optimization, evaluated on ground-robot and handheld platforms. Our setting differs from all of these in sensor and regime: a *single-chip* radar (10–60 points per frame, 2-TX elevation array) on a quadrotor in aggressive racing and attitude-tumbling flight up to  $10 \text{ rad s}^{-1}$ , a regime in which, as Section VII-C shows, the radar measurement model itself degrades in ways invisible on ground vehicles and in benign aerial navigation.

**Single-chip RIO on aerial platforms.** Single-chip radar-inertial fusion has been shown on multirotors, but only for benign flight. Michalczyk et al. [13] use a discrete-time EKF with scan-to-scan point matching and Doppler (a factor-graph reformulation follows in [14]); Girod et al. [15] run a real-time onboard M-estimator smoother on a quadrotor in cities and forests, but cap velocity at  $2 \text{ m s}^{-1}$  and anchor the vertical channel with a barometer. None uses a continuous-time spline, addresses aggressive or tumbling flight, or, as we do, recovers the vertical channel from radar and IMU alone (Section VI-E).

**Sliding-window marginalization and consistency.** That naïve sliding-window marginalization is inconsistent is long established in discrete time [16], where FEJ methods fix linearization points to preserve the observability nullspace. In

continuous time, CLIC [17] and LIO-MARS [18] marginalize spline knots leaving the window by adjacency to the dropped knots; this is correct when the estimated states are all local to the window, but not when a *global* shared IMU bias couples every IMU factor and tempts the adjacency rule to pull in far-edge factors, producing an over-confident prior. Our contribution, narrower than and complementary to FEJ, is a factor-*set* selection rule that exploits B-spline local support to keep the marginal exactly closed despite this shared bias, together with a diagnostic symptom pattern for the inconsistent variant; both are developed in Section IV-E and Section VII-D. In the discrete-time UAV line, the EKF-RIO family of Doer and Trommer [1] defines the state of practice. Absolute root-mean-square error (RMSE) is not comparable across papers: our own quantization ablation (Section VI-D) shows firmware configuration alone changes Doppler resolution by  $12\times$ , so rather than quote their numbers we cross-validate against this line directly: Section VI-E runs our system on the public ICINS-2021 flight datasets of [19] and re-evaluates their EKF estimators under our causal metrics on the same data, and reports the regime-dependent result of porting their filters to our platform.

### III. SYSTEM OVERVIEW

#### A. Problem Formulation

**Estimation problem.** We estimate the continuous-time 6-DoF body trajectory (world-frame position  $\mathbf{p}_w(t) \in \mathbb{R}^3$  and orientation  $\mathbf{R}(t) \in \text{SO}(3)$ ) together with the constant IMU biases  $\mathbf{b}_a, \mathbf{b}_g$ , from asynchronous measurements over the estimation interval (the full trajectory in batch mode, a sliding window online): radar Doppler velocities ( $\sim 11 \text{ Hz}$ ), accelerometer and gyroscope samples ( $1 \text{ kHz}$ ), and a low-rate heading prior. We adopt a *continuous-time* representation because the sensors are asynchronous and run at disparate rates: a trajectory defined for every  $t$  is evaluated at each measurement's exact timestamp, and its derivatives (velocity  $\dot{\mathbf{p}}_w$ , acceleration  $\ddot{\mathbf{p}}_w$ , and body angular velocity  $\boldsymbol{\omega}_b$ ) follow in closed form, so no inter-sample interpolation or rigid-keyframe approximation is required.

**Parameterization.** The trajectory is a B-spline (Section IV-A): the position is a quintic uniform B-spline with control points  $\{\mathbf{P}_i\} \subset \mathbb{R}^3$ , and the orientation a cumulative SO(3) B-spline with *absolute* rotation knots  $\{\mathbf{R}_i\} \subset \text{SO}(3)$  (basalt convention [6]). This renders the otherwise infinite-dimensional trajectory finite-dimensional; the control points and the IMU biases form the estimated state

$$\mathcal{X} = \{\mathbf{P}_i, \mathbf{R}_i, \mathbf{b}_a, \mathbf{b}_g\}. \quad (1)$$

The batch solve additionally frees the single observable extrinsic degree of freedom (the radar pitch); the sliding window fixes it at the measured value (Section IV-A).

**Estimator.** Treating the measurement noise as zero-mean and independent, the maximum-a-posteriori (MAP) estimate is the minimizer of the negative log-posterior, a weighted non-

linear least-squares problem over the measurement residuals, the spline regularizers, and the priors:

$$\hat{\mathcal{X}} = \arg \min_{\mathcal{X}} \sum_k \rho_{\delta}(r_{D,k}^2) + \lambda_a \sum_m \|r_{a,m}\|^2 + \lambda_g \sum_m \|r_{g,m}\|^2 + \lambda_{\psi} \sum_n r_{\psi,n}^2 + E_{\text{reg}}(\mathcal{X}) + E_{\text{prior}}(\mathcal{X}), \quad (2)$$

where  $r_D, r_a, r_g, r_{\psi}$  are the radar-Doppler, accelerometer, gyroscope, and heading residuals of Section IV-B;  $\rho_{\delta}$  is the Huber kernel ( $\delta=1.0$  m/s) that bounds the influence of heavy-tailed Doppler outliers; the  $\lambda_{\bullet}$  are the (rate-adaptive) information weights of Section IV-B;  $E_{\text{reg}}$  collects the position min-snap and orientation smoothness regularizers (Section IV-C); and  $E_{\text{prior}}$  holds the boundary and, online, the marginalization priors (Section IV-E). We minimize (2) with Levenberg-Marquardt (Ceres) using analytic Jacobians. The fixed-lag estimator solves it over a 3 s window and, on each stride, marginalizes the states leaving the window into  $E_{\text{prior}}$  via the Schur complement (Section IV-E), keeping the per-window cost bounded.

### B. Hardware Platform

The system runs on a quadrotor with a TI IWR6843AOPEVM mmWave radar (60–64 GHz FMCW), a Pixhawk flight controller IMU, and a Vicon motion capture (MoCap) system used for ground truth evaluation. The radar is mounted upside-down at a nominal 30° downtilt by computer-aided design (CAD); the as-built angle is a few degrees lower (Section VI), and the extrinsic rotation is  $[180^\circ, 27.5^\circ, 0^\circ]$  (roll, pitch, yaw), translation  $[0.08, 0.02, -0.01]$  m in the body frame. The radar boresight lies along the body forward axis, so the Euler pitch is numerically the physical boresight downtilt. The radar operates at  $\sim 11$  Hz with 0.049 m/s Doppler resolution (best-velocity firmware configuration). Raw IMU rate is  $\sim 993$  Hz; both modalities are used at full rate.

### C. Pipeline

Fig. 1 shows the pipeline. A three-stage sensor-only initialization (P1–P3, detailed in Section IV-D) builds the initial trajectory the nonlinear solve needs, from radar and IMU data alone; the same raw measurements then also enter the solver directly as factors (Fig. 1, dashed). MoCap supplies only the yaw heading prior (a magnetometer surrogate, Section IV-B) and the final RMSE evaluation; P1–P3 and the position/attitude estimate use no external pose reference.

## IV. METHODOLOGY

### A. State Representation

**Position.** The world-frame position  $\mathbf{p}_w(t) \in \mathbb{R}^3$  is a quintic (order-6) uniform B-spline with knot spacing  $\Delta t_p = 5$  ms (slow racing; per-regime grids in Table II). Velocity  $\dot{\mathbf{p}}_w$  and acceleration  $\ddot{\mathbf{p}}_w$  are obtained as analytical derivatives of the spline.




Fig. 1. System pipeline. Solid: initialization flow (P1–P3 seed the optimizer). Dashed: raw sensor data fed as measurement factors. Radar returns first pass a RANSAC ego-velocity prefilter (Section IV-B) that rejects elevation-biased outliers before both P2 and the solver. MoCap supplies only the yaw heading prior (magnetometer surrogate) and RMSE evaluation.

**Orientation.** We use the *cumulative* SO(3) B-spline formulation [6], [20]:

$$\mathbf{R}(t) = \mathbf{R}_{k-3} \prod_{j=k-2}^k \text{Exp}(\tilde{B}_j(t) \text{Log}(\mathbf{R}_{j-1}^{-1} \mathbf{R}_j)), \quad (3)$$

where  $k$  is the active knot index, the  $\mathbf{R}_i \in \text{SO}(3)$  are *absolute* rotation knots (the optimization variables; basalt convention [6]),  $\text{Log}(\mathbf{R}_{j-1}^{-1} \mathbf{R}_j)$  are their relative increments, and  $\tilde{B}_j(t)$  are cumulative cubic basis functions. Knot spacing is  $\Delta t_o = 8$  ms (slow racing; see Table II). Body-frame angular velocity  $\omega_b(t)$  is obtained as a closed-form derivative of (3) directly from `evaluate_lie()` without numerical differentiation.

**Biases and extrinsics.** A constant 6-DoF IMU bias  $[\mathbf{b}_a, \mathbf{b}_g]$  is estimated jointly. The extrinsic pitch  $\delta_{\text{pitch}}$  (1 DoF) is optimized with a soft prior; roll and yaw are fixed (unobservable from Doppler). It is freed in the batch solve but locked at the measured 27.5° for the sliding window (Section VII-C).

### B. Measurement Models

**Radar Doppler with lever arm.** The TI IWR6843 reports positive Doppler for a receding target. With antenna offset  $\mathbf{t}_{bs}$  from the body center of mass (CoM), the predicted Doppler for a point at sensor-frame bearing  $\hat{\mathbf{u}}_s$  is:

$$v_D = -\hat{\mathbf{u}}_b^\top (\mathbf{R}_{bw} \mathbf{v}_w(t) + \omega_b(t) \times \mathbf{t}_{bs}), \quad (4)$$

where  $\hat{\mathbf{u}}_b = \mathbf{R}_{bs} \hat{\mathbf{u}}_s$ ,  $\mathbf{R}_{bs}$  is the rotation from radar sensor to body frame (extrinsic  $[\text{roll}=180^\circ, \text{pitch}=27.5^\circ, \text{yaw}=0^\circ]$ ; measured, Section VI), and  $\mathbf{R}_{bw} = \mathbf{R}_{wb}^\top$ . The leading minus makes the convention self-consistent: a sensor translating toward a static point ( $\hat{\mathbf{u}}_b^\top \mathbf{R}_{bw} \mathbf{v}_w > 0$ ) yields  $v_D < 0$ , an approaching rather than receding target, as required. The residual  $r_k = v_{D,\text{meas}} - v_{D,\text{pred}}$  uses Huber loss with  $\delta = 1.0$  m/s, robust as defense-in-depth to any residual Doppler noise on the surviving inlier returns; the  $0.049 \text{ m s}^{-1}$  Doppler resolution sets a lower bound but does not drive the choice. This  $\delta$  is deliberately large relative to the 0.07–0.08 m/s ego-velocity residuals we ultimately reach (Section VI-E); after the RANSAC prefilter the inliers are already clean, so  $\delta$  caps only the rare survivor rather than shaping the fit.

**RANSAC ego-velocity prefilter (default front-end).** A single-chip radar with only 2 TX antennas has limited elevation diversity, and a minority of returns carry an elevation-correlated Doppler bias that a robust kernel only down-weights. Following reve [1] (Doer’s radar ego-velocity estimator), we therefore prepend a reve-style 3D least-squares

RANSAC ego-velocity estimate to the per-frame Doppler set (inlier threshold 0.15 m/s) and pass only the inlier returns to the spline solver, hard-rejecting the elevation-biased minority before the optimization. The prefilter is seeded (`numpy.default_rng(0)`), so the full pipeline reproduces bit-identically; frames with fewer than five returns bypass it unchanged. On those sparse frames (most of the backflip regime) the in-solver Huber loss is the sole radar outlier protection, so the two stages are complementary. It runs once at frame load, not per window, so it does not enter the per-window timing budget (Section VI). Section VI-E quantifies what it removes: on the public ICINS flights it closes an order-of-magnitude vertical-drift gap, and on our own flights it improves aggressive-racing live position by 20% ( $0.50 \rightarrow 0.40$  m live). It is enabled by default for every result in this paper.

**Dynamics-adaptive measurement weighting.** Our central practical finding is that a single *rate-adaptive trust allocation* across the three measurement streams covers hover to  $10 \text{ rad/s}^{-1}$  with one configuration, each component derived from a measured failure diagnosis (Section VII-C), not tuned ad hoc:

- *Gyroscope: trust more when rotating fast.* Per-sample weight  $\lambda_g^{\text{eff}} = \lambda_g (1 + (|\omega_g|/\omega_g)^p)$  with  $\lambda_g=4$ ,  $\omega_g=4 \text{ rad/s}$ ,  $p=4$ , using the *measured* rate (causal, no spline evaluation). The quartic exponent keeps racing untouched ( $\lambda^{\text{eff}} \approx \lambda_g$  below 3 rad/s) and reaches  $\lambda^{\text{eff}} \approx 160\text{--}500$  during flips, where the gyroscope is the only clean sensor. A quadratic ramp ( $p=2$ ) stiffens the gyro already at 2–5 rad/s (where radar is still informative) and degrades racing roll/yaw.
- *Radar: trust less when rotating fast.* A radar frame integrates over tens of milliseconds; at body rate  $\omega$  the scene smears by  $\omega T_{\text{frame}} (\approx 17^\circ$  at 10 rad/s). Per-frame noise inflation  $\sigma_{\text{eff}}^2 = \sigma_0^2 (1 + (|\omega|/\omega_r)^2)$ ,  $\omega_r=4 \text{ rad/s}$ ,  $|\omega|$  from the warm-start spline at problem-build time; continuous, no hard-gate tuning cliff.
- *Accelerometer: same form*,  $\omega_a=8 \text{ rad/s}$ : during flips the accelerometer carries thrust transients and vibration, its gravity information is nil mid-flip, and it otherwise distorts orientation through the  $\mathbf{R}\text{--}\dot{\mathbf{p}}$  coupling.
- *Per-point signal-to-noise ratio (SNR) weighting.*  $w_i = \text{clip}((I_i/\bar{I})^\alpha, 0.25, 4)$  with frame-median intensity  $\bar{I}$  and  $\alpha=1$  (median-relative, so the global radar weight is unchanged). High-SNR returns carry measurably better Doppler: on aggressive racing this alone improves live orientation from  $3.31^\circ$  to  $2.88^\circ$  at unchanged position.

The weighting trades are made explicit in Section VII-C: on tumbling flight the radar/accelerometer gates exchange velocity accuracy for orientation accuracy along a measured Pareto curve;  $(\omega_r, \omega_a)=(4, 8)$  is its knee.

**Flip-regime Doppler correction.** Under sustained tumbling an additional attitude-correlated Doppler systematic appears, which we absorb with a per-point term  $v_{\text{corr}} = v_{\text{meas}} - b u_{s,z}$  ( $b=-1.5 \text{ m/s}$ , backflips only). Section VII-C shows why  $b$  must be interpreted as a flight-regime error proxy rather than a physical elevation bias of the 2-TX array.

### Accelerometer.

$$\mathbf{r}_a = \mathbf{z}_a - \mathbf{R}_{bw}(\ddot{\mathbf{p}}_w - \mathbf{g}_w) - \mathbf{b}_a, \quad (5)$$

weighted  $\lambda_a = 0.01$ : this 400:1 gyro/accel ratio ( $\lambda_g=4$ , below) is a deliberate design statement that leans position and velocity on radar Doppler plus min-snap, since the accelerometer carries thrust transients and frame vibration that corrupt the  $\mathbf{R}\text{--}\dot{\mathbf{p}}$  coupling (the dynamics-adaptive law of Section IV-B down-weights it further mid-flip).

### Gyroscope.

$$\mathbf{r}_g = \mathbf{z}_g - \omega_b(t) - \mathbf{b}_g, \quad (6)$$

weighted  $\lambda_g = 4.0$  (L2 loss).

### Gravity direction.

$$\mathbf{r}_{\text{tilt}} = \frac{\mathbf{z}_a - \mathbf{b}_a}{\|\mathbf{z}_a - \mathbf{b}_a\|} - \mathbf{R}_{bw}^\top \hat{\mathbf{g}}, \quad (7)$$

active when  $|\|\mathbf{z}_a - \mathbf{b}_a\| - g| < 3 \text{ m/s}^2$ . This provides roll/pitch anchoring during near-hover phases without contaminating dynamic flight.

**Heading prior.** Yaw is unobservable from Doppler alone (all yaw-equivalent trajectories predict identical radial velocities under pure rotation about gravity). We supply a heading reference from a magnetometer (standard on most platforms; here simulated by MoCap yaw). The residual  $r_\psi = \psi_{\text{est}} - \psi_{\text{ref}}$  is added at  $\sim 100 \text{ Hz}$ , with  $\lambda_\psi = 0.6$  by default and raised to 10 on the heading-critical bags (slow racing, backflips; Table III footnote). A weight sweep shows the split is a yaw-observability/position trade. On slow racing, gentle dynamics leave yaw radar-unconstrained, so the weak default ( $\lambda_\psi=0.6$ ) lets yaw drift to  $4.3^\circ$  (orientation  $2.0 \rightarrow 4.2^\circ$ ) while a strong prior costs almost no position. On fast racing, the rich translational excitation already holds yaw near  $3.9^\circ$  at  $\lambda_\psi=0.6$ , so raising it to 10 buys only  $0.6^\circ$  of yaw yet degrades settled position ( $0.31 \rightarrow 0.43 \text{ m}$ ) through the orientation-velocity coupling; we therefore keep it light. On backflips the anchor instead stabilizes the whole attitude through the flip: dropping  $\lambda_\psi$  to 0.6 inflates settled orientation from  $5.0^\circ$  to  $6.8^\circ$  (all three axes degrade), so it stays at 10. The per-regime weight is therefore not monotone in dynamics: gentle (slow racing) and tumbling (backflips) flight need the same strong prior at opposite rates while translation-rich fast racing needs a weak one, so it cannot key off  $|\omega|$ . It is thus the one weight *not* folded into the runtime-adaptive law above, set once per flight type rather than auto-weighted; keying it instead off a yaw-observability proxy such as translational excitation, and adapting it to real magnetometer noise in place of the idealized MoCap heading, is left to future work.

### C. Regularization

**Position: minimum snap.**  $E_{\text{snap}} = \lambda_s \int \|\mathbf{p}^{(4)}(t)\|^2 dt$  with  $\lambda_s = 2 \times 10^{-5}$ .

**Orientation: minimum angular acceleration.** Rather than penalizing angular velocity increments (minimum  $\omega$ ), which opposes every banked turn and the entire backflip maneuver, we penalize the *second* finite difference on SO(3):

$$\mathbf{r}_\alpha = \text{Log}(\mathbf{R}_{i-1}^{-1} \mathbf{R}_i) - \text{Log}(\mathbf{R}_i^{-1} \mathbf{R}_{i+1}), \quad (8)$$

which is zero for constant angular rate. Swept  $\lambda_\alpha \in \{0, 0.001, 0.01, 0.1, 1.0\}$ ; 0.1 was best across all bags (1.0 caused backflip divergence).

#### D. Sensor-Only Initialization (P1–P3)

**P1: Orientation.** A stationary segment before flight is detected automatically; the gyroscope bias is the mean reading over this segment and the accelerometer bias a gravity-aligned scalar correction, both *sensor-only* and using no external attitude reference. Initial roll/pitch are then derived from the stationary accelerometer (gravity direction). Yaw is set to zero (gauge freedom; corrected later by the heading prior). The full orientation trajectory is obtained by forward-integrating debiased gyroscope readings:  $\mathbf{R}(t + \Delta t) = \mathbf{R}(t) \text{Exp}(\mathbf{z}_g - \mathbf{b}_g)\Delta t$ .

**P2: Position.** Per radar frame, weighted least squares estimates the 3D ego-velocity  $\mathbf{v}_{\text{wls}}$  from the RANSAC inlier returns (Section IV-B). Doppler alias unwrapping uses IMU-integrated world-frame velocity to select the correct alias once at frame load (20.9% of points aliased in fast racing, 1.7% in slow racing); the unwrapped Doppler is the measurement consumed by *both* this WLS estimate and the main solver’s radar factors (Eq. 4), with the solver’s own  $v_{\text{max}}$  check kept only as a residual safety net. World-frame velocity is integrated with forward Euler to produce an initial position trajectory.

**P3: Boundary priors.** Position, velocity, orientation, and angular velocity boundary priors are built from the sensor-only state at  $t = 0$ .  $\lambda_{\text{bnd}} = 1000$  for position and velocity.

#### E. Sliding Window with Schur Complement Marginalization

The fixed-lag smoother advances a 3 s window in 0.3 s strides. After each Ceres LM solve, the *stride zone* CPs and orientation knots are marginalized:

$$S = H_{bb} - H_{ab}^\top H_{aa}^{-1} H_{ab}, \quad (9)$$

where  $a$  indexes stride-zone variables and  $b$  indexes boundary variables (last  $N-1$  position CPs, last  $N-1$  orientation knots, and the 6-DoF bias; dimension 30).  $S$  is factored as  $S = LL^\top$  and stored as a `MargPriorFunctor` for the next window.

**Re-centered prior.** Before adding the prior to each window’s problem, the linearization point  $\mathbf{x}_0$  is updated to the current warm-start. This makes the prior contribute only curvature (the Hessian shape), not a gradient pull toward a stale historical estimate. The residual is  $\mathbf{r} = L^\top(\mathbf{x} - \mathbf{x}_0)$  in local coordinates.

**Consistent factor selection (Markov-blanket rule).** Which residuals enter Eq. (9) determines whether the prior is a true marginal. A naïve choice (all residuals touching marginalized *or* boundary variables) pulls in every IMU factor of the window through the shared bias block ( $\sim 25,000$  residual rows), *conditions* on interior states (treating them as perfectly known), and double-counts factors that are re-added in the next window. The resulting prior is overconfident by orders of magnitude and must be deflated by a hand-tuned scale of  $10^{-7}$ – $2 \times 10^{-4}$  depending on flight dynamics (Section VII-D). The shared global bias is the trap: it is what tempts pulling in the far-edge IMU factors. We instead exploit the splines’ local

support of  $N$  knots (quintic position  $N=6$ , cubic orientation  $N=4$ ).

**Proposition 1** (Exact closed spline marginal at a fixed linearization). *At a fixed linearization point, let  $\mathcal{M}$  be the knots leaving the window and  $\mathcal{B}$  the  $N-1$  retained knots immediately forward of  $\mathcal{M}$  plus the global 6-DoF IMU bias (the dim-30 boundary of Section IV-E). Every order- $N$  spline factor has support over  $N$  consecutive knots, and an IMU factor also touches the bias; hence any factor touching  $\mathcal{M}$  has support  $\subseteq \mathcal{M} \cup \mathcal{B}$  with its bias dependence in  $\mathcal{B}$ , while factors reaching further back are already summarized by the incoming prior (re-added here). Selecting exactly the factors touching  $\mathcal{M}$  (plus that prior) therefore forms a Schur complement that conditions on no interior state and shares no factor with the next window: at this linearization it is the exact marginal.*

Including *only* residuals that touch a marginalized variable (rather than those touching marginalized *or* boundary variables) thus yields a closed marginal. Empirically it is also *consistent*: the prior is applied **unscaled** (scale = 1) across all flight regimes, and its Mahalanobis residual at the solution drops from  $10^6$ – $10^{12}$  to  $\mathcal{O}(1)$ , i.e. it agrees with the incoming data.

**Relation to FEJ.** This mechanism is distinct from, and complementary to, first-estimates-Jacobian approaches to marginalization consistency [16]: FEJ fixes the *linearization points* of the prior to preserve the observability nullspace, whereas our rule corrects the *factor set* entering the Schur complement so that the marginal is exactly closed *at a fixed linearization point*. We do not employ FEJ; instead the prior is re-centered at each window’s warm start (curvature retained, gradient pull discarded). We do not claim this re-centering preserves consistency across relinearizations as a theoretical guarantee; rather, in our experiments it left no measurable inconsistency symptom; the NEES analysis of Section VI-F provides the corresponding empirical evidence. Our missions are short (18–26 s), which may be too brief to expose drift a re-centered prior could accumulate over a long deployment; bounding this over extended runs is left to future work.

**Live-edge warm start.** Control points entering the window each stride would otherwise hold stale dead-reckoned initialization values; any CP-level seam is amplified by the min-snap term ( $\propto \Delta/\Delta t_p^4$ , initial cost  $\sim 10^{13}$ ). We rigidly align the entering segment (shift +  $\Delta R$ ) to the last data-supported CP of the previous solution, reducing the initial cost by five orders of magnitude.

**Marginalization cost.** The restricted Gauss–Newton Hessian for Eq. (9) is accumulated by evaluating the affected cost functions directly (tangent-space Jacobians via the manifold plus-Jacobian, Triggs correction for robust losses), avoiding a full solver re-linearization; the marginalization step costs  $\sim 10$  ms per window.

## V. EXPERIMENTAL SETUP

All experiments use the rosbag datasets collected in a Vicon MoCap lab. Table I summarizes the three bags used for evaluation.

TABLE I  
DATASET SUMMARY

| Bag                    | Dur. (s) | $v_{\max}$ (m/s) | $ \omega _{\max}$ (rad/s) | Frames |
|------------------------|----------|------------------|---------------------------|--------|
| Slow racing            | 25.8     | 3.65             | 2.62                      | 258    |
| Fast racing            | 18.0     | 5.66             | 9.14                      | 179    |
| Backflips <sup>*</sup> | 18.1     | $\sim 4.3$ (95%) | $\sim 10.9$ (95%)         | 181    |

<sup>\*</sup>95th percentile; MoCap tracking artifacts dominate the instantaneous peaks. The  $\sim 10$  rad/s headline is this 95th-percentile sustained rate, not an instantaneous maximum.

TABLE II  
PER-REGIME SLIDING-WINDOW CONFIGURATION

| Parameter                            | Slow | Fast | Backflips |
|--------------------------------------|------|------|-----------|
| Position knot $\Delta t_p$ (ms)      | 5    | 40   | 10        |
| Orientation knot $\Delta t_o$ (ms)   | 8    | 16   | 8         |
| LM iteration cap                     | 12   | 400  | 400       |
| Heading prior $\lambda_\psi$         | 10   | 0.6  | 10        |
| Doppler elevation corr. $b$ (m/s)    | —    | —    | -1.5      |
| Position-init anchor $\lambda_{p_0}$ | —    | —    | 0.5       |

Shared across all regimes: 3 s window, 0.3 s stride, unscaled Markov-blanket marginalization prior, the single dynamics-adaptive weighting law (Section IV-B), RANSAC ego-velocity prefilter (0.15 m s<sup>-1</sup> inlier), sensor-only initialization, and frozen extrinsic pitch 27.5°. “—” marks the shared default: uncapped-by-data iterations, and the flight-regime elevation correction / position tether off.

All three regimes run the *same* sliding-window solver, weighting law, and front-end; Table II collects the handful of per-regime knobs that differ, so that every reported number can be reproduced from one place.

**IMU rate:**  $\sim 993$  Hz (full rate, no downsampling for C++ solver). **Evaluation:** Start-anchored rigid alignment to MoCap ground truth: the alignment rotation  $\mathbf{R}_{\text{align}} = \mathbf{R}_{\text{gt},0} \mathbf{R}_{\text{est},0}^\top$  and translation  $\mathbf{t}_{\text{align}} = \mathbf{p}_{\text{gt},0} - \mathbf{R}_{\text{align}} \mathbf{p}_{\text{est},0}$  are fixed at frame 0 (no scale). The final 3 s is trimmed. Velocity RMSE is computed from the analytical first derivative of the position B-spline, not from finite differences of the estimated trajectory. *Position drift* (%) is position RMSE normalised by the total GT path length, the standard odometry metric for a dead-reckoning system with no loop closure (per-segment relative pose error (RPE) in Section VI-B). Start-anchored alignment is harsher than the whole-trajectory SE(3)/Umeyama fit (the EVO/KITTI default) used by most odometry papers, because it cannot retro-fit a constant transform to cancel accumulated drift: it fixes the frame once, at the start, exactly as a deployed estimator must. We adopt it precisely for this causal, deployment-relevant reading, and apply it identically to every system compared. Where we compare against externally published figures (Section VI-E) we additionally report whole-trajectory-aligned numbers so the baselines can be cross-checked against their own papers. For sliding window, *live* (*leading-edge*) RMSE is computed by snapshotting the estimate at  $[t_{\text{end}} - \text{stride}, t_{\text{end}}]$  of each window at 50 Hz before any subsequent refinement. The *settled* RMSE evaluates the same trajectory after all windows have refined each point.

**Solver implementation.** The Ceres LM problem is solved with SPARSE\_NORMAL\_CHOLESKY (SuiteSparse), up to 400 iterations. All three factor types (gyroscope, accelerometer, and radar Doppler) use hand-derived analytic Jacobians based on the basalt cumulative-spline Jacobian struct; Ceres auto-

TABLE III  
SLIDING WINDOW: SETTLED AND LIVE LEADING-EDGE RMSE  
(CONSISTENT UNSCALED PRIOR; PER-WINDOW SOLVE TIME ON A LAPTOP CPU)

| Bag                    | Mode             | Vel (m/s)      | Ori (°)     | Pos (m)      | Drift%      | $t_{\text{win}}$ (s) |
|------------------------|------------------|----------------|-------------|--------------|-------------|----------------------|
| Slow racing            | Settled          | 0.42           | 1.53        | 0.293        | 0.61        | 0.70                 |
|                        | <b>Live edge</b> | <b>0.48</b>    | <b>1.97</b> | <b>0.310</b> | <b>0.64</b> |                      |
| Fast racing            | Settled          | 0.25           | 2.35        | 0.312        | 0.69        | 0.35                 |
|                        | <b>Live edge</b> | <b>0.32</b>    | <b>2.88</b> | <b>0.399</b> | <b>0.88</b> |                      |
| Backflips <sup>†</sup> | Settled          | — <sup>*</sup> | 5.01        | 1.674        | 3.09        | 0.6                  |
|                        | <b>Live edge</b> | <b>2.35</b>    | <b>6.35</b> | <b>1.545</b> | <b>2.85</b> |                      |

All bags use the unscaled consistent marginalization prior (Section IV-E) and the *same* dynamics-adaptive weighting law (Section IV-B). Per-regime parameters are limited to spline grids ( $\Delta t_p = 5/40/10$  ms), the heading weight, and for backflips the position-init anchor ( $\lambda_{\text{pos-init}} = 0.5$ ) and flip-regime Doppler correction  $b = -1.5$  m/s (Section VII-C). Extrinsic pitch locked at 27.5° for all bags. <sup>\*</sup>Settled velocity is not reported for backflips: stride-boundary resampling injects a retrospective-evaluation artifact there, which also leaves backflips settled *position* above live (the one such inversion); the live column is the deployment-relevant figure throughout. Drift %: path lengths 48.1 m / 44.9 m / 54.1 m. Iteration-capped configurations reproduce bit-identically across repeated runs on a fixed platform/BLAS (single-thread reduction); full-convergence configurations vary by <2% (multithreaded reduction order).

matic differentiation was eliminated to avoid Jet-arithmetic overhead over the large set of spline control-point parameters per factor. In the sliding-window solver the linear solve and the marginalization-prior recompute (`compute_prior`, evaluated *every* LM iteration, as distinct from the  $\sim 10$  ms one-time prior *formation* of Section IV-E) are the dominant costs; Jacobian evaluation accounts for roughly 25–30% of per-window wall time.

## VI. RESULTS

### A. Batch Solver and Extrinsic Pitch

We use the full-trajectory C++ batch solve as an offline full-batch reference; offline it reaches 0.20 m / 1.1° on slow racing and 0.64 m / 2.5° on fast (the dense racing  $\Delta t_p$  over-fits fast racing, where the coarser sliding window does better). The radar-to-body pitch is a fixed mount angle, measured at 27–28° with an inclinometer on the radar face (reference surface 0°) and frozen at 27.5°, a few degrees below the 30° CAD design (3D-printed, hand-assembled fixture). We lock it rather than estimate it online because it is only weakly observable from Doppler: the self-calibration is init-dependent and orientation RMSE is pitch-insensitive (Section VII-C).

### B. Sliding Window: Settled vs. Live

Table III reports sliding window results (3 s window, 0.3 s stride, Schur marginalization). Fig. 2 shows trajectory overlays for all three flights; Fig. 3 shows the corresponding error timeseries.

Live-edge orientation degrades 0.4–0.5° relative to settled; this is expected because subsequent windows refine previous estimates. As a complementary per-segment view, the live-edge KITTI translational RPE converges to 1.0–1.7% beyond 20 m on racing (from 5–9% at 5 m segments, where the  $\sim 11$  Hz radar yields few Doppler per segment) and to  $\sim 5.6\%$  on backflips; full curves are released with the code. With the consistent prior, settled velocity is smooth across stride boundaries (the earlier inconsistent prior produced position

TABLE IV  
EVALUATION CONFIGURATION (3 s/0.3 s, USED THROUGHOUT THE PAPER) VS. THE PER-REGIME REAL-TIME POINT (SOLVE TIME BELOW THE STRIDE). REACHING REAL TIME IS ACCURACY-NEUTRAL AND IMPROVES BACKFLIPS. LIVE-EDGE RMSE; SOLVE TIMES CARRY  $\sim 20\%$  RUN-TO-RUN VARIANCE.

| Regime    | Win./str. (s)       | Solve (s) | Live pos (m) | Live ori ( $^\circ$ ) |
|-----------|---------------------|-----------|--------------|-----------------------|
| Fast      | 3.0/0.3 (eval.)     | 0.35      | 0.399        | 2.88                  |
|           | 2.0/0.3 (real-time) | 0.28      | 0.443        | 2.95                  |
| Slow      | 3.0/0.3 (eval.)     | 0.70      | 0.310        | 1.97                  |
|           | 1.5/0.5 (real-time) | 0.40      | 0.301        | 1.98                  |
| Backflips | 3.0/0.3 (eval.)     | 0.60      | 1.545        | 6.35                  |
|           | 2.0/0.5 (real-time) | 0.43      | <b>1.499</b> | <b>5.56</b>           |

jumps that inflated the settled velocity derivative). We report all tables, plots, and the cross-validation of Section VI-E at a single *evaluation configuration* (3 s window, 0.3 s stride), chosen for best accuracy and cross-table consistency, not as a deployment claim: per-window solve time is set by window length and knot density, not the stride, and at 3 s only fast racing approaches the 0.3 s budget on a laptop CPU. Real time is nonetheless within reach: pushing the window and stride to *barely* reach it per regime (Table IV) costs almost nothing. Fast is real-time at the 0.3 s stride with a 2 s window, slow and backflips at a 0.5 s stride (1.5 s and 2 s windows); slow racing and backflips stay within run-to-run noise (backflips even improves,  $6.35 \rightarrow 5.56^\circ$ ), and fast-racing position degrades only modestly, under 1% drift ( $0.88 \rightarrow 0.98\%$ ). Solve times carry  $\sim 20\%$  variance, so the strides are taken with margin.

Deployment asks for more than *barely* real-time. As one factor source in a larger fusion graph (Section VI-F) sharing compute with a parallel visual or lidar pipeline, a tuned, no-margin point on a dedicated CPU is a feasibility demonstration, not a deployment budget. The margin must come from the back-end, and since a banded Schur factorization is numerically unstable at our conditioning, the route is incremental smoothing, not batch re-solve: Girod et al. [15] reach comfortable onboard rates with exactly such an iSAM2 back-end (discrete-time graph), and porting our continuous-time B-spline onto it, where the shared global bias induces a Bayes-tree fill-in, is the deployment path we leave to future work. The dominant cost reductions were: the consistent prior removing the marginalization re-linearization ( $\sim 0.6$  s), coarsening  $\Delta t_p$  from 5 ms to 40 ms for aggressive flight (which *also* drops the LM iteration count from 28 to 11 and slightly improves orientation), and warm-start alignment removing the live-edge seam.

### C. Radar Contribution

Radar is essential. Disabling it (`--no-radar`) collapses the live-edge estimate to IMU dead-reckoning: position drift balloons to 5.9% on slow racing and diverges on fast racing ( $\approx 99.5\%$  of distance), with velocity RMSE rising to 2.4 and 8.0 m/s as double-integrated accelerometer noise accumulates unbounded. The direct Doppler velocity constraints bound this to the Table III figures (0.6–0.9% drift, 0.32–0.48 m/s velocity), one-to-two orders of magnitude lower. Orienta-

tion improves more modestly, from  $8.9^\circ/4.0^\circ$  (IMU-only) to  $2.0^\circ/2.9^\circ$ , the 1 kHz gyroscope already constraining it.

### D. Held-Out Validation

All tuning used three bags (slow/fast racing, backflips); the final configuration was then run *unmodified* on flights never used for any decision. A second, independent fast-racing flight reaches  $0.37\text{ m}/2.81^\circ$  live (vs.  $0.40\text{ m}/2.88^\circ$  on the tuning bag), and a circular trajectory reaches  $0.90\text{ m}/3.77^\circ$  live. Five additional stress-tier bags from an earlier recording day use an older radar firmware with  $12\times$  coarser Doppler quantization ( $0.604$  vs.  $0.049\text{ m/s}$ ), partially unusable accelerometer axes, and per-day time offsets; the default RANSAC prefilter still keeps 46–75% of returns (no starvation), and the system degrades gracefully ( $1.5\text{--}7\text{ m}$  position), with the *old-firmware backflips* orientation reaching  $8.1^\circ$  settled/ $9.1^\circ$  live (better than the  $14.8^\circ$  of its previous, individually tuned best) under the untuned universal configuration. The rate-adaptive weighting thus transfers across both firmware and day. A data-health screening (radar frame-rate continuity, IMU/MoCap coverage) precedes all held-out evaluation, since several early bags contain radar universal asynchronous receiver/transmitter (UART) dropouts; all evaluation windows lie on healthy spans. Each flight condition is a single representative recording; the evaluation's breadth comes from these held-out flights and the cross-platform validation of Section VI-E rather than repeated trials (Section VII-A).

### E. External Cross-Validation Against EKF Baselines

We cross-validate against the EKF-RIO family of Doer and Trommer [1], [19] in both directions on shared data, with their toolbox pinned at a fixed commit and every parameter deviation documented in the repository.

**Our system on their data.** We run our sliding-window solver, with the universal weighting configuration unmodified, on the four public ICINS-2021 indoor flight datasets [19] (TI IWR6843AOP at 10 Hz, Analog Devices ADIS16448 IMU at 409 Hz with onboard barometer, loop-closure-and-scale-corrected VINS pseudo ground truth), using their published extrinsic calibration (locked) and their  $\pm 60^\circ$  elevation pre-gate. Their estimators (ekf-rio, unaided yaw; ekf-yrio, Manhattan-world radar yaw aiding) are re-run from their stock configurations on the same bags, and *both* systems are evaluated with the identical causal protocol of Section V (live estimates, start-anchored alignment, same windows and ground truth). Table V reports the headline metrics (a) and their decomposition (b). On position our *horizontal* drift is competitive ( $0.14\text{--}0.30\%$  of path, Table V(b)); the larger total drift ( $0.6\text{--}1.3\%$  vs. the baselines'  $0.3\text{--}0.6\%$ ) is almost entirely the *vertical* channel ( $0.63\text{--}1.28\%$ ), the single-chip elevation residual the RANSAC prefilter targets (below). Velocity RMSE  $0.07\text{--}0.08\text{ m/s}$  matches theirs; absolute position is  $0.5\text{--}1.9\text{ m}$ , trailing the yaw-aided ekf-yrio. The full 3-DoF orientation column favors our solver ( $0.51\text{--}1.13^\circ$  vs. up to  $10.1^\circ$ ), but this gap is almost entirely *yaw*: our heading is anchored by a pseudo-magnetometer (here a VINS pseudo-GT






Fig. 2. Trajectory overlays for all three flights (columns): MoCap ground truth vs. the SW live edge, X-Y (top) and Y-Z (bottom) projections sharing the Y axis; start ■. Larger fast-racing deviations reflect higher dynamics ( $\omega_{\max} = 9.1 \text{ rad/s}$ ); on backflips the residual  $\sim 6^\circ$  live orientation error attenuates the loop circles while the donut sweep and altitude are followed.




Fig. 3. SW live-edge (deployed) position, velocity, and orientation error vs. time for slow racing (top), fast racing (middle), and backflips (bottom). Fast-racing position error accumulates monotonically with trajectory length; velocity and orientation errors spike at high-speed, high-dynamics sections; backflip errors peak *during* the flips, where the radar measurement model degrades (Section VII-C).

yaw surrogate), whereas ekf-yrio uses sensor-only Manhattan-world yaw and ekf-r10 is unaided. On the heading-*independent* roll and pitch the three systems are comparable ( $0.25\text{--}0.74^\circ$ , Table V(b)), so the orientation column is not an attitude win; removing our heading aiding confirms this: on flight\_1 our unaided yaw drifts to  $9.9^\circ$ , matching unaided ekf-r10 ( $10.1^\circ$ ), while roll/pitch and position are essentially unchanged.

The default RANSAC prefilter (Section IV-B) buys this same-order position, its clearest single demonstration. A single-chip 2-TX array has limited elevation diversity, and a minority of returns carry an elevation-correlated Doppler bias. Doer’s reve rejects these with 3D RANSAC (inlier  $0.15 \text{ m/s}$ ); a Huber-robust least squares only down-weights them. Measured against the ground-truth attitude, the per-frame vertical ego-

velocity bias is  $-0.06$  to  $-0.16 \text{ m/s}$  under a Huber-only front-end but at most  $0.10 \text{ m/s}$  in magnitude under RANSAC. With the prefilter *disabled*, our vertical channel accumulates a large systematic error (whole-trajectory absolute trajectory error (ATE)  $9.6 \text{ m}$  on flight\_1,  $12\text{--}15\%$  vertical drift) that the baselines do not exhibit (theirs stays  $\sim 0.2\text{--}0.3 \text{ m}$ ); the RANSAC prefilter removes it at the source on all four flights, cutting whole-trajectory ATE to  $0.24\text{--}0.76 \text{ m}$  (Table V(b)) and bringing position back into the baselines’ range. Their vertical accuracy is not altitude aiding either: disabling the barometer in ekf-yrio shifts its vertical drift by  $\leq 0.04 \text{ m}$  (stock  $\sigma_{\text{altimeter}} = 5 \text{ m}$  makes it near-negligible), so they too estimate vertical from radar and IMU alone, by the same outlier-rejection mechanism. The single-chip elevation bias is



therefore real but not fundamental: an outlier-rejection gap, closed by the same RANSAC stage the baselines use. Under whole-trajectory Umeyama alignment (Table V(b)) all three systems agree to  $\leq 1$  m, consistent with the baselines' published full-alignment figures (they are faithfully reproduced) and confirming that our default front-end matches them.

**Their filters on our data.** The reverse port is a *sensor/stack-coupling characterization*, not a head-to-head comparison; it is regime-dependent, and we report the measured outcome rather than a verdict. We deliberately keep their *stock* innovation gate and process noise rather than re-tuning them to our platform: the question here is out-of-the-box coupling, and any re-tune would confound it with our own parameter search. Here we port ekf-rio and its multi-radar generalization x-rio [21] (run single-radar); ekf-yrio was the forward-port comparator above. On the slow-racing flight both ekf-rio and x-rio fail: their  $\chi^2(3)$  innovation gate rejects 99.7% of the radar ego-velocity updates (301 of 302), so the filter reverts to IMU dead-reckoning and runs essentially unconstrained ( $> 2$  km of error for both, almost entirely vertical,  $\approx \frac{1}{2}gt^2$ ). On the faster flight ekf-rio instead survives, degrading to 1.4 m (3.1% drift,  $9.0^\circ$ ): its larger translational excitation produces radar ego-velocities that clear the gate (only 64% rejected), enough to constrain the filter. The radar information admitted thus scales with how strongly motion excites the sparse Doppler; the common cause is a stack-level mismatch, not a defect of their method: our best-velocity firmware yields only  $\sim 13$  points per frame (their RANSAC/ $\sigma$ -gates expect hundreds, and ego-velocity estimation fails on roughly half of all scans), there is no hardware radar trigger, and our racing-frame vibration exceeds their process-noise range, making propagation overconfident so the gate rejects the survivors. This is a design-space split (Section VIII), not a defect of either method: their stack runs near sensor rate on dense clouds, ours spends 0.35–0.70 s per window on sparse Doppler, an accuracy-versus-latency trade.

#### F. Output Covariance Consistency (NEES)

A factor source for a larger fusion graph must report *honest* covariance. Each window we extract the live-edge joint covariance of  $(\mathbf{v}_w, \mathbf{R})$  from the converged problem (`ceres::Covariance`, active support control points, tangent space, marginalization prior included) and compute the NEES against MoCap; for a consistent estimator each 3-DoF block is  $\chi_3^2$  with *mean* 3. We report the per-window *mean* after a ground-truth-quality mask (the dropout/glitch mask used on backflips, applied per channel; the finite-difference MoCap velocity reference is additionally discarded where a near-duplicate timestamp makes it return an implausible speed). On the racing flights the *orientation* NEES mean is 1.6 (slow) and 3.2 (fast), within or just below the  $\chi_3^2$  95% interval ( $\approx [2.0, 3.7]$ ): orientation covariance is calibrated on fast racing and slightly conservative on slow. Velocity is calibrated where the finite-difference reference is clean (slow racing, mean 2.7); on the faster and tumbling flights that reference is itself too noisy ( $\sim 1$  m/s) for a powered check. These are *calibration checks*, not powered hypothesis tests: consecutive

TABLE V  
CROSS-VALIDATION ON THE PUBLIC ICINS-2021 FLIGHT DATASETS [19], ALL SYSTEMS UNDER OUR CAUSAL PROTOCOL (LIVE ESTIMATE, START-ANCHORED ALIGNMENT). (A) HEADLINE METRICS; HEADING AIDING DIFFERS BY DESIGN (SEE TEXT). (B) DECOMPOSITION, MAKING THE HONEST STRUCTURE VISIBLE AT A GLANCE: *roll/pitch* IS THE HEADING-INDEPENDENT ATTITUDE RMSE (PER-AXIS RMS, YAW EXCLUDED), SO THE LARGE ORIENTATION GAP IN (A) IS YAW, NOT ATTITUDE; *our drift* SPLITS OUR POSITION ERROR INTO HORIZONTAL AND VERTICAL, BOTH SUB-1.3% UNDER THE DEFAULT RANSAC PREFILTER (SECTION IV-B; THE FOOTNOTE GIVES THE PREFILTER-DISABLED FIGURES FOR CONTRAST); *whole-traj.* ATE RE-EVALUATES EVERY SYSTEM UNDER THE STANDARD EVO/KITTI UMEYAMA ALIGNMENT (NO SCALE) FOR CROSS-CHECKING AGAINST PUBLISHED FIGURES.

| (a) drift [% path] / velocity [m/s] / orientation [°] |                          |                          |  |          |                          |  |      |  |
|-------------------------------------------------------|--------------------------|--------------------------|--|----------|--------------------------|--|------|--|
| Flight (path)                                         | ekf-rio                  |                          |  | ekf-yrio |                          |  | Ours |  |
| 1 (142 m)                                             | 1.0 / 0.20 / 10.1        | <b>0.3 / 0.08 / 0.70</b> |  |          | 0.6 / 0.08 / <b>0.51</b> |  |      |  |
| 2 (45 m)                                              | 0.6 / 0.08 / 1.59        | <b>0.4 / 0.08 / 0.91</b> |  |          | 1.1 / 0.08 / 0.94        |  |      |  |
| 3 (147 m)                                             | <b>0.3 / 0.08 / 1.07</b> | <b>0.3 / 0.08 / 1.09</b> |  |          | 1.3 / <b>0.07 / 0.52</b> |  |      |  |
| 4 (78 m)                                              | 0.5 / 0.09 / 3.67        | <b>0.3 / 0.07 / 1.41</b> |  |          | 1.2 / 0.07 / <b>1.13</b> |  |      |  |

| (b) Decomposition (systems: ekf-rio / ekf-yrio / Ours) |                |      |      |               |      |               |      |      |
|--------------------------------------------------------|----------------|------|------|---------------|------|---------------|------|------|
| Fl.                                                    | Roll/pitch [°] |      |      | Our drift [%] |      | Whole ATE [m] |      |      |
|                                                        | rio            | yrio | Ours | hor.          | ver. | rio           | yrio | Ours |
| 1                                                      | 0.24           | 0.25 | 0.25 | 0.16          | 0.63 | 0.87          | 0.23 | 0.46 |
| 2                                                      | 0.57           | 0.59 | 0.55 | 0.14          | 1.13 | 0.13          | 0.10 | 0.24 |
| 3                                                      | 0.25           | 0.26 | 0.27 | 0.30          | 1.28 | 0.36          | 0.26 | 0.76 |
| 4                                                      | 0.74           | 0.74 | 0.74 | 0.26          | 1.15 | 0.30          | 0.15 | 0.46 |

Prefilter *disabled* (Huber-only): our vertical drift 12–15%; whole-traj. ATE 9.57/2.88/10.86/5.48 m (flights 1–4).

windows are temporally correlated, so the  $\chi^2$  expectation holds only heuristically. On backflips the orientation covariance is overconfident (NEES mean 38,  $\sim 3.5\times$  in standard deviation, unmodeled flip-time error), and only 6 of 50 windows survive the ground-truth mask (Section V), too few for a consistency statement, so we report it as a caveat. Deployment recipe: inflate the emitted covariance per regime ( $1\times$  gentle,  $3\times$  aggressive), or calibrate online from innovations.

#### G. Ablation Studies

Table VI summarizes key design choices.

**Key observations from Table VI.** Full-rate IMU (1 kHz) improves position by  $2.9\times$  over 200 Hz, motivating the tight bias priors in `solver_cpp.yaml`. Lever-arm correction helps racing (8–11% position) but has negligible effect on backflips: at  $\Delta t_o = 0.8$  ms the dense spline absorbs the  $|\omega \times \mathbf{t}_{bs}| \approx 1$  m/s term through orientation DoF. Schur marginalization reduces fast-racing settled position by 24% ( $0.412$  m  $\rightarrow$   $0.312$  m); for slow racing the optimal prior scale is  $10^{-7}$  (near-zero), so both variants give equivalent results; the benefit is captured in Table VII instead. Window duration has little effect on either racing regime: slow stays at live 0.30–0.34 m across 1.5–5 s, and fast gives essentially equal orientation at 2 s and 3 s (settled  $2.31^\circ$  vs  $2.35^\circ$ ) with 3 s modestly better on position (live drift 0.98 vs 0.88%). We keep 3 s as the default for that small position gain, trading latency. We lock pitch at its measured value for all sliding-window runs; even in the batch ablation, locking is marginally better than freeing it (0.157 vs. 0.169 m on slow racing), consistent with the pitch being only weakly observable (Section VII-C).

TABLE VI  
SELECTED ABLATION RESULTS (BATCH SOLVER, SLOW\_RACING UNLESS NOTED; SW ROWS USE THE DEFAULT 0.3 S STRIDE)

| Configuration                                                             | Pos (m)              | Ori (°)            |
|---------------------------------------------------------------------------|----------------------|--------------------|
| <i>IMU rate (slow_racing batch, no lever arm<sup>†</sup>)</i>             |                      |                    |
| 200 Hz                                                                    | 0.525                | 1.64               |
| <b>1000 Hz (default)</b>                                                  | <b>0.183</b>         | <b>1.02</b>        |
| <i>Lever-arm correction <math>\omega \times t_{bs}</math> (batch)</i>     |                      |                    |
| No correction — slow racing                                               | 0.183                | 1.02               |
| — fast racing                                                             | 0.862                | 2.73               |
| — backflips                                                               | 1.814                | 8.64               |
| <b>With correction — slow</b>                                             | <b>0.169</b>         | <b>1.02</b>        |
| — fast                                                                    | <b>0.768</b>         | <b>2.64</b>        |
| — backflips                                                               | <b>1.814</b>         | <b>8.44</b>        |
| <i>Orientation regularization (backflips batch)</i>                       |                      |                    |
| Min- $\omega$ ( $\lambda_\omega=0.001$ )                                  | 1.816                | 8.62               |
| <b>Min-<math>\alpha</math>, <math>\lambda_\alpha=0.1</math> (default)</b> | <b>1.814</b>         | <b>8.44</b>        |
| Min- $\alpha$ , $\lambda_\alpha=1.0$                                      | 2.069                | 83.8               |
| <i>SW window dur. (slow_racing SW, settled / live edge)</i>               |                      |                    |
| 1.5 s                                                                     | 0.292 / 0.299        | 1.51 / 2.01        |
| <b>3.0 s (default)</b>                                                    | <b>0.293 / 0.310</b> | <b>1.53 / 1.97</b> |
| 5.0 s                                                                     | 0.313 / 0.341        | 1.49 / 2.03        |
| <i>SW window dur. (fast_racing SW, settled / live edge)</i>               |                      |                    |
| 2.0 s                                                                     | 0.380 / 0.443        | 2.31 / 2.95        |
| <b>3.0 s (default)</b>                                                    | <b>0.312 / 0.399</b> | <b>2.35 / 2.88</b> |
| <i>Marginalization (fast_racing SW, settled)<sup>‡</sup></i>              |                      |                    |
| None (boundary prior only)                                                | 0.412                | 2.71               |
| <b>Schur complement</b>                                                   | <b>0.312</b>         | <b>2.35</b>        |
| <i>Extrinsic optimization (slow_racing batch)<sup>§</sup></i>             |                      |                    |
| <b>Pitch <math>\delta</math> optimized (default)</b>                      | <b>0.169</b>         | <b>1.02</b>        |
| Pitch $\delta$ locked                                                     | 0.157                | 1.02               |

Batch rows (IMU-rate, lever-arm, ori-reg, extrinsic) isolate one design choice at a fixed pre-RANSAC operating point; the RANSAC default shifts batch absolutes (slow batch 0.169→0.198 m / 1.02→1.1°, fast 0.768→0.64 m, the Section VI batch figures) without changing orderings. The SW rows (window duration, marginalization) are under the RANSAC default. All runs use the sensor-only init (Section IV-D); the bias seed shifts batch absolutes comparably, again without changing orderings. <sup>†</sup>Lever arm disabled for isolation; see Lever-arm rows for the current default (0.169 m/1.02°). <sup>‡</sup>Marginalization shown for *fast\_racing*; on *slow\_racing* the near-zero optimal prior makes Schur and boundary-only nearly equivalent. <sup>§</sup>C++ solver supports only 1-DoF pitch offset (roll/yaw unobservable from Doppler; free 3-angle optimization tested in Python solver showed 6–8° orientation runaway; see Section VII-E).

## VII. LIMITATIONS AND NEGATIVE RESULTS

### A. Statistical Scope

Each flight regime is evaluated on a *single* representative recording, so the headline numbers carry no run-to-run error bars. The held-out flights (Section VI-D), the cross-platform validation (Section VI-E), and the deterministic, bit-identical re-runs broaden the evidence, but they do not substitute for repeated trials: re-runs add reproducibility, not statistical power. Reporting mean  $\pm$  std over several recordings per regime is the natural strengthening, and is left to future data collection.

### B. Heading Sensor Requirement

Yaw is a gauge freedom for Doppler: rotating the entire trajectory about the world gravity axis leaves all radial velocity measurements unchanged, so absolute heading cannot be recovered from radar and IMU alone. A magnetometer resolves this and is standard on virtually all flight platforms; here MoCap yaw serves that role. Because that MoCap yaw is

noise- and bias-free (unlike a real magnetometer), the heading-aided orientation numbers are in principle optimistic. We bound this by re-running with a magnetometer surrogate on the heading prior (a slowly-varying 2° heading drift, a half-period sine that the start-anchored alignment cannot cancel, plus 1.5° white noise, seeded): live orientation degrades by at most 0.3° (fast racing 2.88 → 3.17°; slow 1.97 → 1.88°, backflips 6.35 → 6.34° unchanged). The 1 kHz gyroscope constrains yaw *rate* so the slow drift is resisted, the white noise averages out, and backflip orientation error is roll/pitch- not yaw-dominated, so the headline is robust to heading-sensor imperfection; the heading-*independent* roll/pitch decomposition (Table V(b)) and the unaided-yaw check (Section VI-E) further isolate attitude from this idealization. This is a sensor-architecture requirement, not an open research problem.

### C. Extreme Dynamics: Backflips

**Batch.** The batch solver handles backflips only with  $\Delta t_o=0.8$  ms (the *batch* density, 10× denser than the 8 ms sliding-window grid of Table II; 22,630 knots over 18 s) and locked extrinsics (free pitch drifts to +18° due to dense under-constrained orientation DoF). At this density the batch orientation is *bistable* (detailed below): under small timestamp perturbations it oscillates between  $\sim 8^\circ$  and  $\sim 25^\circ$  regimes, so we do not report a single batch figure and take the sliding-window result as the trustworthy one.

**Sliding window.** The fixed-lag smoother initially appeared fundamentally unsuited to backflips: with the inconsistent prior and a (silently inactive) position anchor, the best result was 2.53 m/10.8° settled. A chain of three fixes revised this verdict. (i) The consistent unscaled prior (Section IV-E) restores inter-window information transfer. (ii) A soft position anchor to the dead-reckoned initialization ( $\lambda_{\text{pos-init}}=0.5$  per CP) bounds the position random walk that sparse, rate-inflated radar cannot prevent within a 3 s window (without it position diverges to 8 m; on racing bags the same anchor is *harmful*: a flip-regime rescue, not a general ingredient). (iii) The dynamics-adaptive weighting law (Section IV-B), in particular rate-adaptive gyroscope trust, which on this bag alone is worth  $\sim 3^\circ$  of orientation at zero position cost, plus the flip-regime Doppler correction  $b=-1.5$  m/s. Together these reach **1.67 m/5.0° settled and 1.55 m/6.4° live at  $\sim 0.6$  s per window** (Table III; Fig. 2). The default RANSAC prefilter (Section IV-B) is neutral here (backflip frames are already sparse and mostly fall below its five-return floor, so they bypass it), and the small change from the Huber front-end (1.51 → 1.55 m live) is within run-to-run noise.

*Why  $\Delta t_o=8$  ms in the window.* At the batch density  $\Delta t_o=0.8$  ms a 3 s window has  $\sim 11,250$  orientation DoF against  $\sim 9,000$  gyroscope constraints (1.26 DoF/constraint); the per-window Jacobian is rank-deficient and the solver freezes near its initialization. At 8 ms the window is well-conditioned (62.5 Hz model bandwidth comfortably covers the 10–20 Hz flip dynamics). The same density is what makes the *batch* problem bistable (the  $\sim 8^\circ/\sim 25^\circ$  oscillation noted above): the final cost *anti-correlates* with accuracy (the underconstrained problem overfits the wrong global orientation), so single-run comparisons at this density are not trustworthy.

*Attributing the z-axis failure, and what b really is.* A visual trajectory check (numeric RMSE conceals it) showed the SW estimate sinking in altitude. Disabling the position anchor exposes the raw behaviour: a near-constant  $-0.57$  m/s world- $z$  drift ( $-8$  m over the flight). Sweeping the per-point correction  $v_{\text{corr}} = v - bu_{s,z}$  produces a linear family of  $z$ -error curves and large negative  $b$  flattens the drift, so we initially interpreted  $b$  as the antenna-fixed elevation bias of the 2-TX array. A joint calibration experiment *disproves* this: (i) with pitch *locked* at its measured value, level-flight racing degrades catastrophically under any explicit  $b$  (fast-racing live position  $0.39 \rightarrow 3.2 \rightarrow 6.6$  m for  $b=0, -0.5, -1.0$ ), whereas a true antenna-fixed bias would be tolerated; and (ii) on backflips the optimum grows monotonically past the WLS-measured  $-0.5$  m/s (through  $-1.5$  to  $-2.0$ ) into physically implausible territory. We conclude  $b$  is a *flight-regime error proxy*: it absorbs an attitude-correlated Doppler systematic during flips (most likely intra-frame motion distortion) that happens to project like an elevation term, and we cap it at  $b=-1.5$  rather than chase the monotone trend without a mechanism. Relatedly, extrinsic pitch is a *plateau*: over a  $3 \times 3$  grid over (locked pitch,  $b$ ) the RMSE is indistinguishable across  $26.5$ – $28.5^\circ$  on backflips, and locking at  $27.5^\circ$  versus online self-calibration is RMSE-neutral on racing. We freeze it rather than estimate it because free pitch is only weakly observable and *init-dependent* (the converged value tracks the seed:  $25^\circ \rightarrow 27^\circ$ ,  $30^\circ \rightarrow 36^\circ$  on both racing bags), and fix it at the physically measured as-built angle ( $27.5^\circ$ , inclinometer; Section VI), a few degrees below the  $30^\circ$  CAD design (the 3D-printed, loosely-fastened fixture accounts for the difference).

*Remaining gap.* The residual  $\sim 6^\circ$  live orientation error during flips attenuates the reconstructed loop circles relative to ground truth. This is not a spline-bandwidth limit (Section VII-E: denser knots do not help); the residual lever is the radar measurement model, whose intra-frame Doppler smearing during flips is best addressed by *per-chirp timestamps* rather than knot density, which we leave to future work.

#### D. Prior Scale Sensitivity as a Symptom of Inconsistent Marginalization

Our initial marginalization implementation selected all residuals touching marginalized or boundary variables, which, through the shared bias block, baked every IMU factor of the window into the Schur complement, conditioned on interior states, and double-counted factors re-added in the next window (Section IV-E). The resulting overconfident prior had to be deflated by a hand-tuned scale, and the required scale varied by orders of magnitude across flight regimes with a harmful intermediate region. We document this behaviour because the symptom pattern (no universal prior weight, non-monotone sensitivity, robust losses on the prior acting backwards) is a useful *diagnostic for marginalization inconsistency* in sliding-window estimators generally. Table VII shows the sweep measured with the inconsistent prior.

Slow racing is best served by a near-zero prior ( $10^{-7}$ ): with gentle dynamics, each 3 s window is self-sufficient and additional inter-window coupling is harmful. Fast racing instead needs the prior ( $2 \times 10^{-4}$ ) for inter-window position

TABLE VII  
PRIOR SCALE SENSITIVITY (SW, LIVE-EDGE RMSE)

| scale                                      | Slow racing  |                  | Fast racing  |                  |
|--------------------------------------------|--------------|------------------|--------------|------------------|
|                                            | Pos (m)      | Ori ( $^\circ$ ) | Pos (m)      | Ori ( $^\circ$ ) |
| $2 \times 10^{-4}$                         | 0.543        | 2.16             | <b>0.825</b> | <b>3.64</b>      |
| $10^{-5}$                                  | —            | —                | 0.942        | 3.58             |
| $10^{-6}$                                  | 0.442        | 2.21*            | 0.951        | 3.57             |
| <b><math>10^{-7}</math> (slow default)</b> | <b>0.338</b> | <b>2.08</b>      | 1.279        | 3.57             |
| $10^{-8}$                                  | 0.336        | 2.08             | 1.344        | 3.57             |

\*Orientation local maximum:  $2.21^\circ > 2.16^\circ$  ( $2e-4$ ) and  $2.08^\circ$  ( $1e-7$ ). —: scale not evaluated for this bag.

Measured under the inconsistent-prior diagnostic configuration (predates the RANSAC default); the sensitivity *pattern*, not the absolutes, is the result.

continuity; without it (scale =  $10^{-7}$ ) fast-racing live position degrades from 0.825 m to 1.279 m. The two bags thus want opposite scales, and the non-monotone slow-racing orientation maximum at  $10^{-6}$  (Table VII) shows no single scale is universally optimal.

The information asymmetry in the inconsistent  $S$  (orientation eigenvalues  $\sim 5.5 \times 10^{14}$ , position  $\sim 10^4$ ) meant no single scalar could address both modalities, and eigenvalue clipping and Cauchy robustification also failed (Cauchy down-weights the bag that needs the prior most).

**Resolution.** With the Markov-blanket factor-selection rule of Section IV-E, the prior is a true marginal: a single *unscaled* prior (scale = 1) is optimal or near-optimal on all three bags simultaneously, the prior's Mahalanobis residual at each solution is  $\mathcal{O}(1)$ , and the entire per-mission tuning table becomes obsolete. The fix also improves accuracy: fast-racing settled position improves from 0.727 m (tuned inconsistent prior) to 0.701 m (unscaled consistent prior). Both figures isolate the prior-consistency effect at a fixed pre-RANSAC operating point; the RANSAC default brings the same metric to 0.312 m (Table III). Stride-boundary position jumps in the settled trajectory also disappear (settled velocity RMSE  $0.89 \rightarrow 0.21$  m/s on slow racing).

#### E. Other Negative Results

**Adaptive knot density.** We tested whether denser orientation knots reduce the residual backflip attitude error; they do not. At  $\Delta t_o = 8$  ms the orientation spline already represents the MoCap attitude to  $1.28^\circ$  RMSE *during* the flips, tracking the gyro stream to 0.295 rad/s RMS (flat across  $\Delta t_o \in \{16, 8, 4\}$  ms, at the rate-independent vibration floor 0.17 rad/s). The gap is therefore *not* a spline-bandwidth limit: the spline can represent the motion but the optimizer does not rotate to it (open-loop gyro dead-reckoning,  $7.8^\circ$ , already beats the  $9.2^\circ$  this fusion reaches at the diagnostic configuration, before the rate-adaptive gyro trust of Section VII-C brings it to the  $6.4^\circ$  headline). The lever is measurement weighting (Section VII-C), not knot density, *disproving* the adaptive-knot-density hypothesis and obviating the non-uniform-basis extension.

**IMU preintegration.** Replacing per-sample gyro/accel factors with Forster TRO-2017 [22] preintegrated factors at 100 Hz reduces constraint density from 8:1 (samples per knot) to 1:1, degrading orientation RMSE from  $1.02^\circ$  to  $6.40^\circ$  and

position RMSE from 0.169 m to 0.288 m on slow racing. The smaller Jacobian iterates faster, but the accuracy loss is unacceptable.

### VIII. CONCLUSION

*Can a single-chip mmWave radar plus IMU provide usable 6-DoF state estimation?* Yes, as a *velocity and orientation* source: at the live leading edge,  $0.48 \text{ m s}^{-1}/2.0^\circ$  on gentle racing,  $0.32 \text{ m s}^{-1}/2.9^\circ$  on aggressive racing (0.35–0.70 s per window on a laptop), confirmed on held-out flights and with an approximately calibrated (NEES) orientation covariance on racing. Absolute position is not observable from Doppler+IMU and drifts with distance (0.6–0.9% on racing, 2.9% under tumbling); yaw requires a heading sensor. Within those structural limits, *good results took three things*: (1) a *consistent* fixed-lag smoother, with a Markov-blanket factor-selection rule that makes the marginalization a true marginal, eliminating regime-dependent prior tuning whose symptom pattern we document as a reusable diagnostic; (2) a *dynamics-adaptive trust allocation* across the measurement streams (rate-adaptive gyroscope weight, rate-inflated radar/accelerometer noise, per-point SNR weighting), one configuration from hover to  $10 \text{ rad s}^{-1}$ , with the velocity–orientation trade made explicit as a measured Pareto curve; and (3) *characterized sensor systematics*, including the finding that the apparent elevation bias under tumbling is a flight-regime error proxy, not an antenna constant, and that the extrinsic pitch is weakly observable, best fixed at its measured value rather than estimated online. Notably, the largest accuracy gains came from diagnosis-led re-weighting of existing measurements, not from added model capacity: a representation-level study showed the orientation spline already tracks  $10 \text{ rad s}^{-1}$  flips at the gyroscope noise floor, disproving our own adaptive-knot-density hypothesis without implementing the non-uniform-knot extension.

Open limitations: (1) loop-scale position geometry through flips is bounded by during-flip radar data quality; (2) ground truth itself degrades during flips (MoCap occlusion), bounding what can be measured there; and (3) cross-validation against the EKF-RIO family (Section VI-E) reveals a design-space split. With our default RANSAC front-end our solver reaches the same order of magnitude on position on their public flights (0.5–1.9 m, trailing the yaw-aided baseline), matches their roll/pitch, and leads on heading-aided orientation; the prefilter closes the single-chip elevation-bias gap. In reverse, their stock filters, starved by our sparse  $\sim 13$ -point clouds and racing-frame vibration, reject up to 99.7% of radar updates and dead-reckon on slow racing. Each stack is engineered for its own regime.

*Future directions* follow from these results: an incremental, iSAM2-style back-end to turn the demonstrated real-time feasibility (Section VI-B) into onboard margins; per-chirp radar timestamps to target the residual during-flip orientation error our negative results place in the measurement model, not the spline (Section VII-C); repeated recordings per regime for powered statistics (Section VII-A); and auto-adapting the heading weight from a yaw-observability proxy rather than setting it per regime (Section IV-B).

### REFERENCES

- [1] C. Doer and G. F. Trommer, “An EKF based approach to radar inertial odometry,” in *IEEE International Conference on Multisensor Fusion and Integration (MFI)*, 2020, pp. 152–159.
- [2] A. Kramer, C. Stahoviak, A. Santamaria-Navarro, A.-a. Aghamohammadi, and C. Heckman, “Radar-inertial ego-velocity estimation for visually degraded environments,” in *IEEE International Conference on Robotics and Automation (ICRA)*, 2020, pp. 5739–5746.
- [3] K. Harlow, H. Jang, T. D. Barfoot, A. Kim, and C. Heckman, “A new wave in robotics: Survey on recent mmWave radar applications in robotics,” *IEEE Transactions on Robotics*, vol. 40, pp. 4544–4560, 2024.
- [4] D. Kellner, M. Barjenbruch, J. Klappstein, J. Dickmann, and K. Dietmayer, “Instantaneous ego-motion estimation using doppler radar,” in *16th Int. IEEE Conf. Intelligent Transportation Systems (ITSC)*, 2013, pp. 869–874.
- [5] P. Furgale, J. Rehder, and R. Siegwart, “Unified temporal and spatial calibration for multi-sensor systems,” in *IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*, 2013, pp. 1280–1286.
- [6] V. Usenko, N. Demmel, D. Schubert, J. Stückler, and D. Cremers, “Visual-inertial mapping with non-linear factor recovery,” *IEEE Robotics and Automation Letters (RA-L)*, vol. 5, no. 2, pp. 422–429, 2020.
- [7] S. Leutenegger, S. Lynen, M. Bosse, R. Siegwart, and P. Furgale, “Keyframe-based visual-inertial odometry using nonlinear optimization,” *International Journal of Robotics Research (IJRR)*, vol. 34, no. 3, pp. 314–334, 2015.
- [8] T. Qin, P. Li, and S. Shen, “VINS-Mono: A robust and versatile monocular visual-inertial state estimator,” *IEEE Transactions on Robotics*, vol. 34, no. 4, pp. 1004–1020, 2018.
- [9] Y. Z. Ng, B. Choi, R. Tan, and L. Heng, “Continuous-time radar-inertial odometry for automotive radars,” in *IEEE/RSJ Int. Conf. Intelligent Robots and Systems (IROS)*, 2021, pp. 323–330, arXiv:2201.02437.
- [10] K. Burnett, A. P. Schoellig, and T. D. Barfoot, “Continuous-time radar-inertial and lidar-inertial odometry using a gaussian process motion prior,” *IEEE Transactions on Robotics*, vol. 41, pp. 1059–1076, 2025, arXiv:2402.06174.
- [11] V.-J. Štirtonja, L. Petrović, J. Peršić, I. Marković, and I. Petrović, “Radar-inertial odometry with online spatio-temporal calibration via continuous-time IMU modeling,” 2026, arXiv:2603.19958.
- [12] Q. Huang, Y. Liang, Z. Qiao, S. Shen, and H. Yin, “Less is more: Physical-enhanced radar-inertial odometry,” in *IEEE Int. Conf. Robotics and Automation (ICRA)*, 2024, pp. 15 966–15 972, arXiv:2402.02200.
- [13] J. Michalczyk, R. Jung, and S. Weiss, “Tightly-coupled EKF-based radar-inertial odometry,” in *IEEE/RSJ Int. Conf. Intelligent Robots and Systems (IROS)*, 2022, pp. 12 336–12 343.
- [14] J. Michalczyk, J. Quell, F. Steidle, M. G. Müller, and S. Weiss, “Tightly-coupled factor graph formulation for radar-inertial odometry,” in *IEEE/RSJ Int. Conf. Intelligent Robots and Systems (IROS)*, 2024, pp. 3364–3370.
- [15] R. Girod, M. Hauswirth, P. Pfreundschuh, M. Biasio, and R. Siegwart, “A robust baro-radar-inertial odometry M-estimator for multicopter navigation in cities and forests,” in *IEEE Int. Conf. Multisensor Fusion and Integration (MFI)*, 2024, pp. 1–8, arXiv:2408.05764.
- [16] T.-C. Dong-Si and A. I. Mourikis, “Consistency analysis for sliding-window visual odometry,” in *IEEE Int. Conf. Robotics and Automation (ICRA)*, 2012, pp. 5202–5209.
- [17] J. Lv, X. Lang, J. Xu, M. Wang, Y. Liu, and X. Zuo, “Continuous-time fixed-lag smoothing for lidar-inertial-camera SLAM,” *IEEE/ASME Transactions on Mechatronics*, vol. 28, no. 4, pp. 2259–2270, 2023, arXiv:2302.07456.
- [18] J. Quenzel and S. Behnke, “LIO-MARS: Non-uniform continuous-time trajectories for real-time lidar-inertial-odometry,” 2025, arXiv:2511.13985.
- [19] C. Doer and G. F. Trommer, “Yaw aided radar inertial odometry using manhattan world assumptions,” in *28th Saint Petersburg Int. Conf. Integrated Navigation Systems (ICINS)*, 2021, pp. 1–9.
- [20] C. Sommer, V. Usenko, D. Schubert, N. Demmel, and D. Cremers, “Efficient derivative computation for cumulative B-splines on Lie groups,” in *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2020, pp. 11 148–11 156.
- [21] C. Doer and G. F. Trommer, “x-RIO: Radar inertial odometry with multiple radar sensors and yaw aiding,” *Gyroscope and Navigation*, vol. 12, no. 4, pp. 329–339, 2021.
- [22] C. Forster, L. Carlone, F. Dellaert, and D. Scaramuzza, “On-manifold preintegration for real-time visual-inertial odometry,” *IEEE Transactions on Robotics*, vol. 33, no. 1, pp. 1–21, 2017.